

Working with Data in Julia

John Myles White

January 14, 2013

```
Pkg.add("DataFrames")  
Pkg.add("RDatasets")  
using DataFrames  
using RDatasets
```

How do we cope with missing data?

```
v = [x1, x2, x3, x4, x5]
```

```
mean(v)
```

```
v = [0.5, 0.6, 0.7, 0.8, 0.9]
```

```
mean(v)
```

```
v = [0.5, 0.6, 0.7, NA, 0.9]
```

```
mean(v)
```

The NA type:

- ▶ Represents a missing value
 - ▶ Like NULL in some systems
- ▶ Poisons other values
 - ▶ Like NaN for floating point numbers

```
1 + NA
```

```
1 > NA
```

```
isna(NA)
```


- ▶ `dataArray{T}` extends `Array{T}`
- ▶ `dataArray{T}` can store T or NA

```
dv = DataArray([1, 2, 3])  
dv[1] = NA  
join(dv, "::")
```

Convenience constructors:

- ▶ `datazeros()`
- ▶ `dataones()`
- ▶ `datafalses()`
- ▶ `datatrues()`
- ▶ `dataeye()`
- ▶ `datadiagm()`

```
dm = dataeye(4)  
svd(dm)
```

Convenience converters:

- ▶ `dataint()`
- ▶ `datafloat()`
- ▶ `databool()`

```
dataint([1.0, 2.0, 3.0])
```

```
databool([1, 0, 1])
```

How do we store data efficiently?

`PooledDataArray{T}` compresses `DataArray{T}`


```
pda = PooledDataArray(["String 1",  
                        "String 2",  
                        "String 2",  
                        "String 1"])  
  
pda[1] = NA
```

```
levels(pda)  
dump(pda)
```

How do we cope with heterogeneous data?

Name	Height	Weight	Gender
John Smith	73.0	NA	Male
Jane Doe	68.0	130	Female

```
df = DataFrame()
df["Name"] = DataVector["John Smith", "Jane Doe"]
df["Height"] = DataVector[73.0, 68.0]
df["Weight"] = DataVector[NA, 130]
df["Gender"] = DataVector["Male", "Female"]
df
```

- ▶ A DataFrame is a list of DataVector's
- ▶ DataFrame's allow mixed indexing:
 - ▶ Columns by number
 - ▶ Columns by name
 - ▶ Rows + Columns by number + number
 - ▶ Rows + Columns by number + name

```
df = DataFrame()  
df["A"] = dataones(5)  
df[2] = datazeros(Int, 5)  
df  
df[1, 1]  
df[1, "A"]
```

```
movies = read_table(joinpath(julia_pkgdir(),  
                              "DataFrames",  
                              "test",  
                              "data",  
                              "movies.csv"))
```



```
head(movies)
tail(movies)
movies[1:2, 1:5]
```

```
iris = data("datasets", "iris")  
colnames(iris)  
coltypes(iris)
```

```
colnames!(iris, colnames(iris))  
clean_colnames!(iris)  
colnames(iris)
```

```
size(iris)
```

```
nrow(iris)
```

```
ncol(iris)
```

```
vcats(iris, iris)
hcats(iris, iris)

rbinds(iris, iris)
cbinds(iris, iris)
```

```
iris[1, 1] = NA  
head(iris)  
complete_cases(iris)
```

```
iris[complete_cases(iris), :]  
complete_cases!(iris)  
iris
```

```
cut([1, 2, 3, 4, 5, 6], [2, 3])
```



```
xtabs([1, 2, 3, 3, 4, 2, 3])
```

```
any(duplicated(iris))  
any(duplicated(rbind(iris, iris)))
```

```
new_iris = rbind(iris, iris)
drop_duplicates!(new_iris)
new_iris
```

```
vector(iris["Species"])  
vector(iris["Species"], Any)
```

```
matrix(iris)
matrix(iris[, 2:3])
matrix(iris[, 1:3])
matrix(iris[, 1:3], Any)
```

```
with(iris, :(Petal_Length .* Petal_Width))  
within!(iris,  
         :(Petal_Area = Petal_Length .* Petal_Width))  
head(iris)
```

Database operations on DataFrames:

- ▶ CRUD
- ▶ subset
- ▶ merge
- ▶ groupby

```
df1 = DataFrame({"a" => [1, 2, 3],  
                  "b" => ["America", "Europe",  
                           "Africa"]})  
df2 = DataFrame({"a" => [1, 2, 4],  
                  "c" => ["New World", "Old World",  
                           "New World"]})  
merge(df1, df2)
```



```
merge(df1, df2, "a", "inner")  
merge(df1, df2, "a", "left")  
merge(df1, df2, "a", "right")  
merge(df1, df2, "a", "outer")
```

The Split-Apply-Combine Strategy:

- ▶ Segment data into groups
- ▶ Apply a function to each group independently
- ▶ Combine results into a single DataFrame

```
by(movies, "year", nrow)  
subset(movies, :(year .== 1893))
```

```
by(movies, ["Action", "year"], nrow)  
subset(movies, :(year .== 1893))
```

Let's do some simple machine learning

```
using Clustering
iris = data("datasets", "iris")
k_means(matrix(iris[:, 2:5]), 3)
iris["Cluster"] =
  k_means(matrix(iris[:, 2:5]), 3).assignments
by(iris, ["Cluster", "Species"], nrow)
```

```
using knn, Resampling
iris = data("datasets", "iris")
train, test = splitrandom(iris, 0.80)
test["Guess"] = knn(matrix(train[:, 2:5]),
                      matrix(test[:, 2:5]),
                      vector(train[:, 6]),
                      10)
by(test, ["Species", "Guess"], nrow)
mean(test["Guess"] .== test["Species"])
```

Simple exercises:

- ▶ Create some DataArray's
 - ▶ A DataArray containing the first five primes
 - ▶ A DataArray containing only NA's
 - ▶ A DataArray containing DataArray's
- ▶ Create some DataFrame's
 - ▶ A 4x3 DataFrame w/:
 - ▶ 4 String's
 - ▶ 4 Int's
 - ▶ 4 ComplexPair's

Advanced exercises:

- ▶ For each species in `iris`:
 - ▶ Find the median petal area
 - ▶ Find the variance of the petal area
 - ▶ Find the centroid of the sepal and petal dimensions